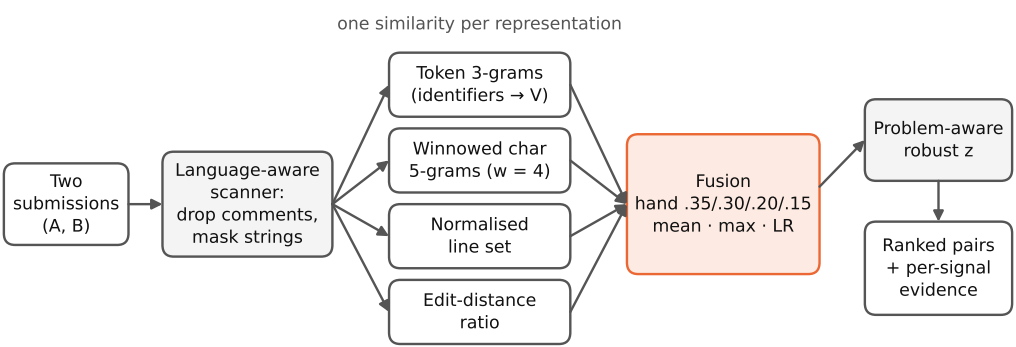
one similarity per representation
Two submissions (A, B)
Language-aware scanner: drop comments, mask strings
Token 3-grams (identifiers → V)
Winnowed char 5-grams (w = 4)
Normalised line set
Edit-distance ratio
Fusion
hand .35/.30/.20/.15
mean · max · LR
Problem-aware robust z
Ranked pairs + per-signal evidence